


# Convolutional Kitchen Sinks for Transcription Factor Binding Site Prediction

Alyssa Morrow\*, **Vaishaal Shankar\***  
Anthony Joseph, Benjamin Recht, Nir Yosef

# Transcription Factor?

A protein that binds to DNA to regulate gene expression



Transcription Factors (TF) attach to specific “binding sites” on DNA

# Why are binding sites important?

- Locating TF binding sites can greatly help understanding of gene regulation
- Presence/Absence of specific TF binding sites can help differentiate between benign and malignant mutations
- Genetic variation in TF binding sites linked to common diseases

Can we predict TF binding sites with data?

# Sequence Classification

Lets set up the machine learning problem

CGTAGTAAC...CGGTAGGA  
AGTCCTAG...GGATCAAT  
ATCGAATCG...GTAGGTACA  
⋮  
ATCGAATCG...GTAGGTACA  
AGATTAGCT...AATGAAAGT

DNA Sequences\*

Bound	Unbound	Unbound
Unbound	Bound	Unbound
Unbound	Bound	Unbound
⋮	⋮	⋮
Bound	Bound	Unbound
Unbound	Unbound	Unbound

ATF2

EGR1

CEBPB

Transcription Factors

Training Cell Type: GM12878

# Sequence Classification

Lets set up the machine learning problem

ACGTGAT...AATGAAAT

TATAGGATC...GGGTACCT

GTCAACG...GTAGGTACA

⋮

GATTTCGG...GTAGGTACA

CGATAC...GGAAGTAAAGTA

*DNA Sequences\**

*Training Cell Type: MCF7*

Unbound Bound Unbound

Unbound Bound Bound

Unbound Bound Bound

⋮

Bound Unbound Unbound

Unbound Unbound Bound

ATF2

EGR1

CEBPB

*Transcription Factors*

# Sequence Classification

Lets set up the machine learning problem

GGACCAT...TACAGAAT

AACGTTAC...AGGACGAT

TAGATTAC...AGATTACA

⋮

AAATTCAG...AAATACCC

GCGATAAC...GTTACGGG

DNA Sequences\*

Unbound

Unbound

Unbound

Unbound

Bound

Unbound

Unbound

Bound

Bound

⋮

⋮

⋮

Unbound

Unbound

Unbound

Unbound

Unbound

Bound

ATF2

EGR1

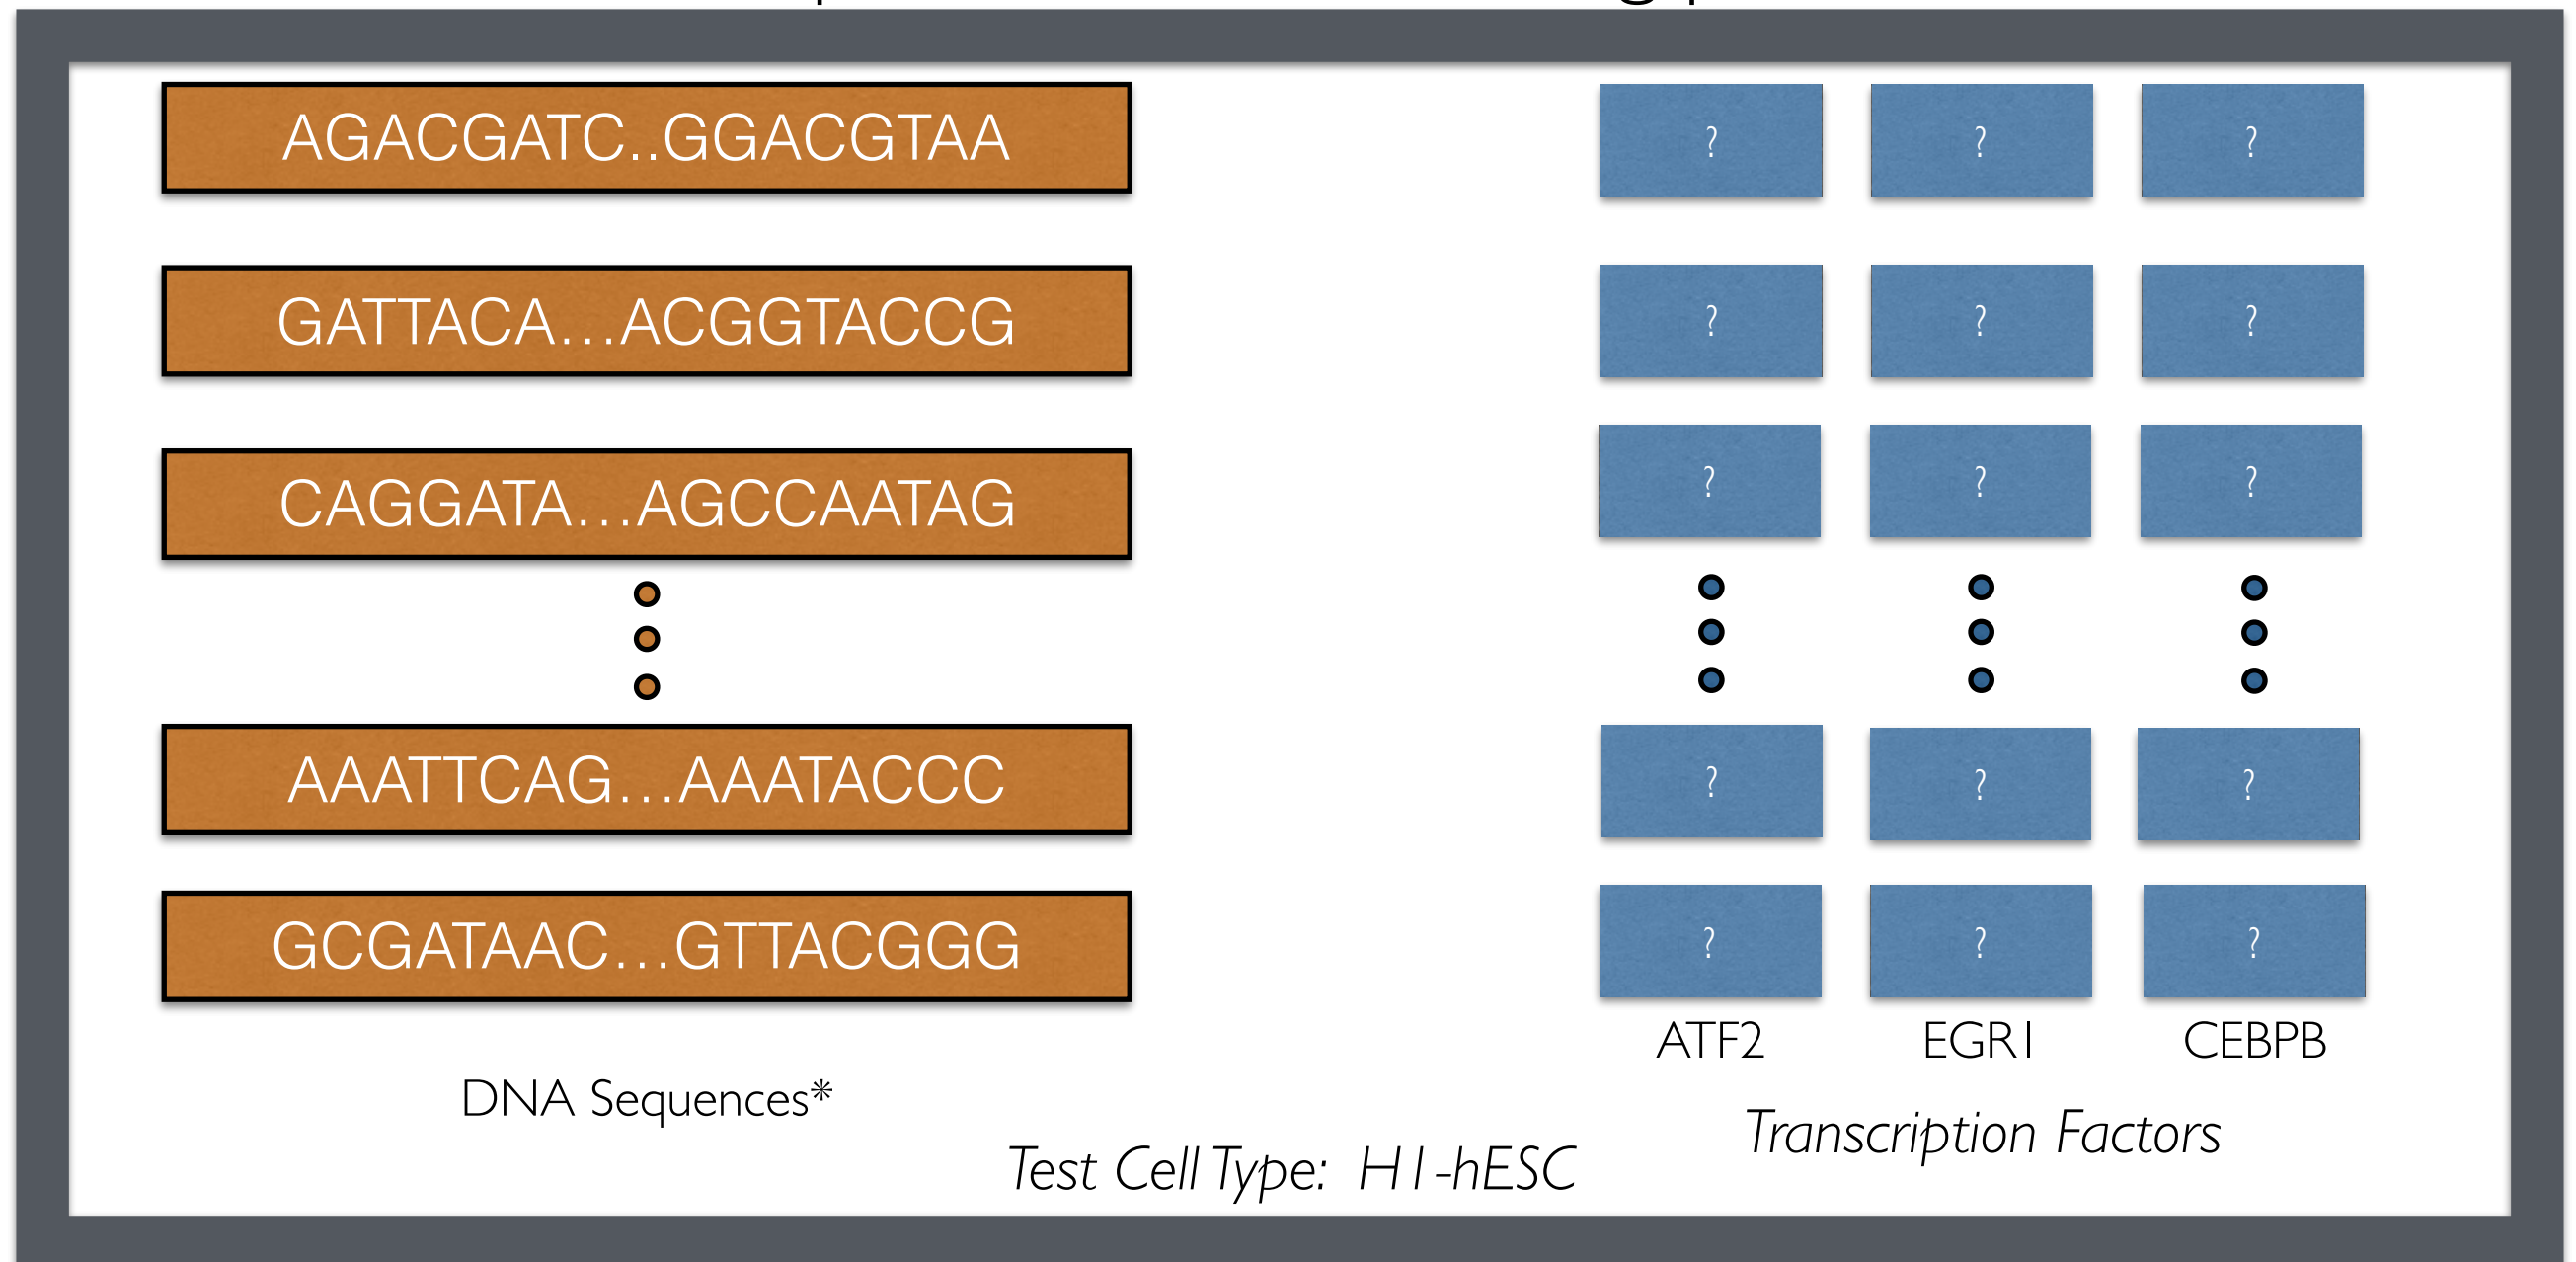
CEBPB

Training Cell Type: HCT116

Transcription Factors

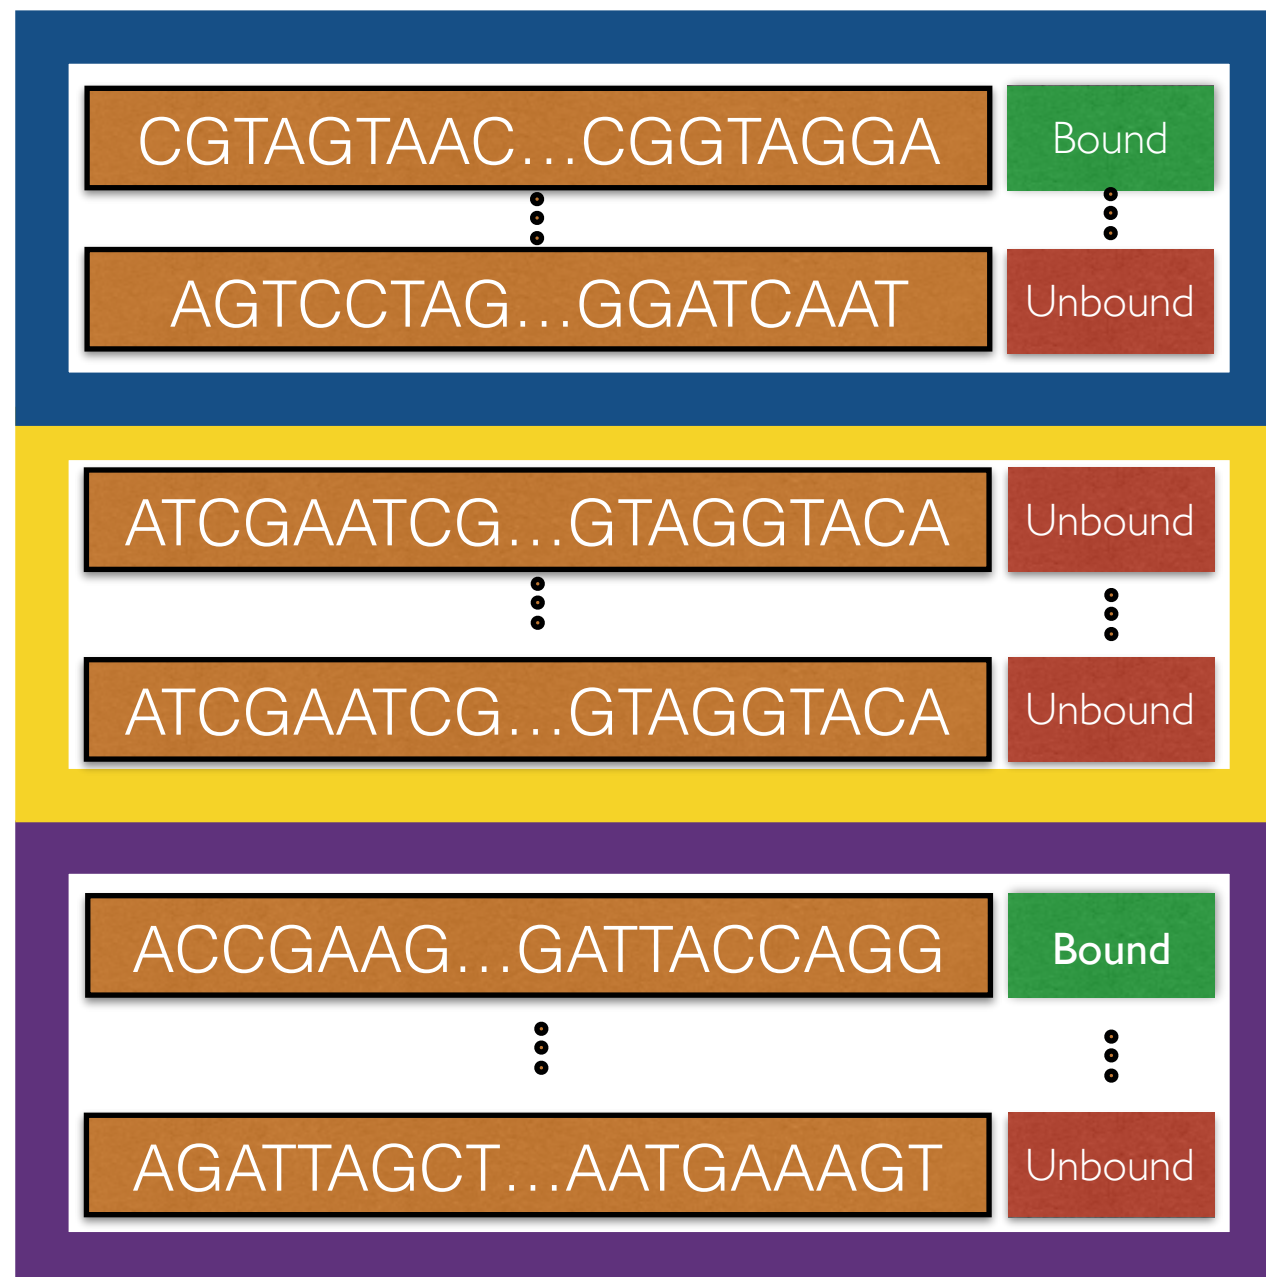
# Sequence Classification

Lets set up the machine learning problem






# Train/Test on single TF



**ATF2 Training Data**  
(Multiple cell types)



**ATF2 Test Data**  
(One Foreign Cell Type)

# We could use a Convolutional Neural Network

- Deep Convolutional Neural Networks (CNNs) are state of the art for many classification tasks
- Recently successful for TF binding site prediction
  - DeepBind
  - Basett
  - DeepSea

# Challenges with Deep Nets

- Scales to large dataset sizes
  - Slow to train
  - Difficult to parallelize
- Design decisions
  - How many layers?
  - How wide?
  - Learning rate?

Can we try something faster & simpler while  
preserving predictive performance?

Yes.

# Kernels

- Well understood theoretically
- Previously successful for biological sequence classification
  - Spectrum Kernel
  - Gapped k-mer Kernel
- Single point of design: kernel function

# How to Design Kernel?

- $k(x,y)$  large for “similar” sequences
  - sequences  $x,y$  that both bind/don't bind to a particular transcription factor
- $k(x,y)$  small for “dissimilar” sequences
- Binding sites are 6-30 base pairs long
  - Local sequence comparisons
  - Aggregation of local comparisons

# A Convolutional String Kernel

ATCGAA

$x_0$

$x$

GGATCG

$y_0$

$y$

$$k(x, y) =$$



# A Convolutional String Kernel

ATCGAA

$x_0$

$x$

GGATCG

$y_0$

$y$

$$k(x, y) = \exp(-\gamma \mathbb{H}^2(x_0, y_0))$$

# A Convolutional String Kernel

ATCGAA

$x_0$

$x$

GGATCG

$y_1$

$y$

$$k(x, y) = \exp(-\gamma \mathbb{H}^2(x_0, y_0)) + \exp(-\gamma \mathbb{H}^2(x_0, y_1))$$

# A Convolutional String Kernel

ATCGAA

$x_0$

$x$

GGATCG

$y_2$

$y$

$$\begin{aligned} k(x, y) = & \exp(-\gamma \mathbb{H}^2(x_0, y_0)) \\ & + \exp(-\gamma \mathbb{H}^2(x_0, y_1)) \\ & + \exp(-\gamma \mathbb{H}^2(x_0, y_2)) \end{aligned}$$

# A Convolutional String Kernel

ATCGAA  
 $x_0$

$x$

GGATCG  
 $y_2$

$y$

$$\begin{aligned} k(x, y) = & \exp(-\gamma \mathbb{H}^2(x_0, y_0)) \\ & + \exp(-\gamma \mathbb{H}^2(x_0, y_1)) \\ & + \exp(-\gamma \mathbb{H}^2(x_0, y_2)) \\ & \vdots \end{aligned}$$

# A Convolutional String Kernel

ATCGAA

$x_i$

$x$

GGATCG

$y_j$

$y$

$$k(x, y) = \sum_{i=0}^{d-n} \sum_{j=0}^{d-n} \exp(-\gamma \mathbb{H}^2(x_i, y_j))$$

# A Convolutional String Kernel

ATCGAA

$x_i$

$x$

GGATCG

$y_j$

$y$

$$k(x, y) = \sum_{i=0}^{d-n} \sum_{j=0}^{d-n} \exp(-\gamma \mathbb{H}^2(x_i, y_j))$$

# What's wrong?

Still have to solve learning problem!

Let  $K$  be the kernel matrix derived by kernel function and training data

Let  $y$  be the training labels

Solve for  $w$

$$Kw = y$$

Matrix Inverse!

Unfortunately  $K$  grows quadratically with dataset size

# Kernel Trick (in reverse)

$$K = \Phi\Phi^T$$

Due to representer theorem it suffices to compute:

$$\min_z \|\Phi z - y\|_2^2$$

Least Squares!

Unfortunately  $\Phi$  is potentially infinite dimensional



The convolutional kernel can be efficiently approximated!

# Definitions

100,000

$N$

Number of training data points

101

$d$

Length of each string in training set

8

$n$

Local comparison window length

0.1

$\gamma$

Kernel bandwidth

4096

$M$

Approximation dimension

# More Definitions

Let  $W$  be a  $M \times n$  matrix

Such that  $W_{ij} \sim N(0, \gamma)$

Let  $b \in \mathbb{R}^M$

Such that  $\forall_i b_i \sim U(0, 2\pi)$

Let  $\hat{\phi}$  be a map from  $\mathbb{R}^n \rightarrow \mathbb{R}^M$

Such that  $\hat{\phi}(x) = \cos(Wx + b)$

# Random Approximation

$$K = \Phi \Phi^T$$

$$k(x, y) = \sum_{i=0}^{d-n} \sum_{j=0}^{d-n} \exp(-\gamma \mathbb{H}^2(x_i, y_j))$$

$$k(x, y) \approx \sum_{i=0}^{d-n} \sum_{j=0}^{d-n} \hat{\phi}(x_i)^T \hat{\phi}(y_j)$$

$$k(x, y) \approx \left( \sum_{i=0}^{d-n} \hat{\phi}(x_i) \right)^T \left( \sum_{j=0}^{d-n} \hat{\phi}(y_j) \right)$$

$$K \approx \hat{\Phi} \hat{\Phi}^T$$

# Kernel Trick (in reverse)

$$K \approx \hat{\Phi} \hat{\Phi}^T$$

Due to representer theorem it suffices to compute:

$$\min_z \|\hat{\Phi} z - y\|_2^2$$

$\hat{\Phi}$  is  $N \times M$  dimensional

If  $M$  is small this problem is easy!

# Putting it Together

1. Compute random map for each training point
2. Solve least squares system with output features
3. Compute *same* random map for each test point
4. Use least squares model to predict labels for test points

# Comparison: DeepBind


- DeepBind is state of art deep learning method for predicting binding sites from DNA sequence
- Trained and tested on DeepBind's datasets:
  - Top 101 bp positive sequences from ENCODE
  - Synthetically generated negative sequences

# Training Time

- Deep bind model takes 2 GPU hours to train for single TF
  - Over 8 hours with cross validation to find optimal hyper parameters
- Our convolutional kitchen sink model takes 10 minutes on a GPU to train for a single TF
  - Less than 20 minutes with cross validation to find optimal hyper parameters




# Accuracy Results



# More Test Data

- Tested TF EGR1 on new sequences preprocessed from ENCODE:
  - Tested on 4 different (foreign) cell types
  - 101 bp positive sequences from ENCODE
  - True negative sequences from genome

# More Accuracy Results



# Conclusion & Future Work

- Fast simple learning algorithm is competitive for transcription factor binding site prediction
- Using epigenetic data such as Dnase hypersensitivity and Histone modification will further help binding site prediction
- Binding site detection (as opposed to recognition) more difficult problem
- Adaptive kernel learning



vaishaal@berkeley.edu

FIN



akmorrow@berkeley.edu

# References

- Leslie, C., Eskin, E., Weston, J., Noble, W. The Spectrum Kernel: A String Kernel for SVM Protein Classification. Proceedings of the Pacific Symposium on Biocomputing, 2002. pp. 564–575.
- Ghandi M. et al. (2014) Enhanced regulatory sequence prediction using gapped k-mer features. PLoS Comput. Biol., 10, e1003711.
- ENCODE Project Consortium et al. The ENCODE (Encyclopedia of DNA elements) project. Science, 306(5696):636–640, 2004.
- Babak Alipanahi, Andrew DeLong, Matthew T Weirauch, and Brendan J Frey. Predicting the sequence specificities of DNA-and RNA-binding proteins by deep learning. Nature biotechnology, 2015.
- Ali Rahimi and Benjamin Recht. Weighted sums of random kitchen sinks: Replacing minimization with randomization in learning. In Advances in Neural Information Processing Systems, pages 1313–1320, 2009.
- Julien Mairal, Piotr Koniusz, Zaid Harchaoui, and Cordelia Schmid. Convolutional kernel networks. In Advances in Neural Information Processing Systems, pages 2627–2635, 2014.
- Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In Advances in Neural Information Processing Systems, pages 1177–1184, 2007.